

Vision 3D artificielle

Session 1: Projective geometry, camera matrix, panorama

Pascal Monasse
monasse@imagine.enpc.fr

IMAGINE, École des Ponts ParisTech



Contents

Pinhole camera model

Projective geometry

Homographies

Panorama

Internal calibration

Optimization techniques

Conclusion

Practical session

Contents

Pinhole camera model

Projective geometry

Homographies

Panorama



Internal calibration

Optimization techniques

Conclusion

Practical session

The “pinhole” camera model



Model

Projection (Source: Wikipedia)

The “pinhole” camera (French: sténopé):

- ▶ Ideal model with an aperture reduced to a single point.
- ▶ No account for blur of out of focus objects, nor for the lens geometric distortion.

Central projection in camera coordinate frame

- ▶ Rays from C are the same: $\vec{C}\vec{X} = \lambda\vec{C}\vec{X}$
- ▶ In the camera coordinate frame $CXYZ$:

$$\begin{pmatrix} x \\ y \\ f \end{pmatrix} = \lambda \begin{pmatrix} X \\ Y \\ Z \end{pmatrix}$$

- ▶ Thus $\lambda = f/Z$ and

$$\begin{pmatrix} x \\ y \end{pmatrix} = f \begin{pmatrix} X/Z \\ Y/Z \end{pmatrix}$$

- ▶ In pixel coordinates:

$$\begin{pmatrix} u \\ v \end{pmatrix} = \begin{pmatrix} \alpha x + c_x \\ \alpha y + c_y \end{pmatrix} = \begin{pmatrix} (\alpha f)X/Z + c_x \\ (\alpha f)Y/Z + c_y \end{pmatrix}$$

- ▶ αf : focal length *in pixels*, (c_x, c_y) : position of principal point P in pixels.

Contents

Pinhole camera model

Projective geometry

Homographies

Panorama

Internal calibration

Optimization techniques

Conclusion

Practical session

Projective plane

- ▶ We identify two points of \mathbb{R}^3 on the same ray from the origin through the equivalence relation:

$$\mathcal{R} : \mathbf{x} \mathcal{R} \mathbf{y} \Leftrightarrow \exists \lambda \neq 0 : \mathbf{x} = \lambda \mathbf{y}$$

- ▶ Projective plane: $\mathbb{P}^2 = (\mathbb{R}^3 \setminus O) / \mathcal{R}$
- ▶ Point $(x \ y \ z) = (x/z \ y/z \ 1)$ if $z \neq 0$.
- ▶ The point $(x/\epsilon \ y/\epsilon \ 1) = (x \ y \ \epsilon)$ is a point “far away” in the direction of the line of slope y/x . The limit value $(x \ y \ 0)$ is the infinite point in this direction.
- ▶ Given a plane of \mathbb{R}^3 through O , of equation $aX + bY + cZ = 0$. It corresponds to a line in \mathbb{P}^2 represented in homogeneous coordinates by $(a \ b \ c)$. Its equation is:

$$(a \ b \ c) (X \ Y \ Z)^T = 0.$$

Projective plane

- ▶ Line through points \mathbf{x}_1 and \mathbf{x}_2 :

$$\ell = \mathbf{x}_1 \times \mathbf{x}_2 \text{ since } (\mathbf{x}_1 \times \mathbf{x}_2)^T \mathbf{x}_i = |\mathbf{x}_1 \quad \mathbf{x}_2 \quad \mathbf{x}_i| = 0$$

- ▶ Intersection of two lines ℓ_1 and ℓ_2 :

$$\mathbf{x} = \ell_1 \times \ell_2 \text{ since } \ell_i^T (\ell_1 \times \ell_2) = |\ell_i \quad \ell_1 \quad \ell_2| = 0$$

- ▶ Line at infinity:

$$\ell_\infty = \begin{pmatrix} 0 \\ 0 \\ 1 \end{pmatrix} \text{ since } \ell_\infty^T \begin{pmatrix} x \\ y \\ 0 \end{pmatrix} = 0$$

- ▶ Intersection of two “parallel” lines:

$$\begin{pmatrix} a \\ b \\ c_1 \end{pmatrix} \times \begin{pmatrix} a \\ b \\ c_2 \end{pmatrix} = (c_2 - c_1) \begin{pmatrix} b \\ -a \\ 0 \end{pmatrix} \in \ell_\infty$$

Calibration matrix

- ▶ Let us get back to the projection equation:

$$\begin{pmatrix} u \\ v \end{pmatrix} = \begin{pmatrix} fX/Z + c_x \\ fY/Z + c_y \end{pmatrix} = \frac{1}{Z} \begin{pmatrix} fX + c_x Z \\ fY + c_y Z \end{pmatrix}$$

(replacing αf by f)

- ▶ We rewrite:

$$Z \begin{pmatrix} u \\ v \\ 1 \end{pmatrix} := \mathbf{x} = \begin{pmatrix} f & c_x \\ & f & c_y \\ & & 1 \end{pmatrix} \begin{pmatrix} X \\ Y \\ Z \end{pmatrix}$$

- ▶ The 3D point being expressed in another orthonormal coordinate frame:

$$\mathbf{x} = \begin{pmatrix} f & c_x \\ & f & c_y \\ & & 1 \end{pmatrix} (R \quad T) \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix}$$

Calibration matrix


- ▶ The (internal) **calibration matrix** (3×3) is:

$$K = \begin{pmatrix} f & & c_x \\ & f & c_y \\ & & 1 \end{pmatrix}$$

- ▶ The **projection matrix** (3×4) is:

$$P = K \begin{pmatrix} R & T \end{pmatrix}$$

- ▶ If pixels are trapezoids, we can generalize K :



$$K = \begin{pmatrix} f_x & s & c_x \\ & f_y & c_y \\ & & 1 \end{pmatrix} \text{ (with } s = -f_x \cotan \theta \text{)}$$

Theorem

Let P be a 3×4 matrix whose left 3×3 sub-matrix is invertible. There is a unique decomposition $P = K \begin{pmatrix} R & T \end{pmatrix}$.

Proof: Gram-Schmidt on rows of left sub-matrix of P starting from last row (RQ decomposition), then $T = K^{-1}P_4$.

Contents

Pinhole camera model

Projective geometry

Homographies

Panorama

Internal calibration

Optimization techniques

Conclusion

Practical session

Homographies

Let us see what happens when we take two pictures in the following particular cases:

1. **Rotation around the optical center** (and maybe change of internal parameters).

$$\mathbf{x}' = K' R K^{-1} \mathbf{x} := H \mathbf{x}$$

2. **The world is flat.** We observe the plane $Z = 0$:

$$\mathbf{x}' = K \begin{pmatrix} R_1 & R_2 & R_3 & T \end{pmatrix} \begin{pmatrix} X \\ Y \\ 0 \\ 1 \end{pmatrix} = K \begin{pmatrix} R_1 & R_2 & T \end{pmatrix} \mathbf{x} := H \mathbf{x}$$

In both cases, we deal with a 3×3 invertible matrix H , a homography.

Property: a homography preserves alignment. If $\mathbf{x}_1, \mathbf{x}_2, \mathbf{x}_3$ are aligned, then

$$|H\mathbf{x}_1 \quad H\mathbf{x}_2 \quad H\mathbf{x}_3| = |H| |\mathbf{x}_1 \quad \mathbf{x}_2 \quad \mathbf{x}_3| = 0$$

Contents

Pinhole camera model

Projective geometry

Homographies

Panorama

Internal calibration

Optimization techniques

Conclusion

Practical session

Panorama construction

- ▶ We stitch together images by correcting homographies. This assumes that the scene is flat or that we are rotating the camera.
- ▶ Homography estimation:

$$\lambda \mathbf{x}' = H\mathbf{x} \Rightarrow \mathbf{x}' \times (H\mathbf{x}) = 0,$$

which amounts to 2 independent linear equations per correspondence $(\mathbf{x}, \mathbf{x}')$.

- ▶ 4 correspondences are enough to estimate H (but more can be used to estimate through mean squares minimization).



Panorama from 14 photos

Algebraic error minimization

- ▶ $\mathbf{x}_i' \times (H\mathbf{x}_i) = 0$ is a system of three linear equations in H .
- ▶ We gather the unknown coefficients of H in a vector of 9 rows

$$\mathbf{h} = (H_{11} \quad H_{12} \quad \dots \quad H_{33})^T$$

- ▶ We write the equations as $A_i \mathbf{h} = 0$ with

$$A_i = \begin{pmatrix} x_i & y_i & 1 & 0 & 0 & 0 & -x_i'x_i & -x_i'y_i & -x_i' \\ 0 & 0 & 0 & x_i & y_i & 1 & -y_i'x_i & -y_i'y_i & -y_i' \\ -x_iy_i' & -y_iy_i' & -y_i' & x_i'x_i & x_i'y_i & x_i' & 0 & 0 & 0 \end{pmatrix}$$

- ▶ We can discard the third line and stack the different A_i in A .
- ▶ \mathbf{h} is a vector of the kernel of A (8×9 matrix)
- ▶ We can also suppose $H_{3,3} = h_9 = 1$ and solve

$$A_{:,1:8} \mathbf{h}_{1:8} = -A_{:,9}$$


Geometric error

- ▶ When we have more than 4 correspondences, we minimize the algebraic error

$$\epsilon = \sum_i \|\mathbf{x}'_i \times (H\mathbf{x}_i)\|^2,$$

but it has no geometric meaning.


- ▶ A more significant error is geometric:



- ▶ Either $d'^2 = d(\mathbf{x}', H\mathbf{x})^2$ (transfer error) or
 $d^2 + d'^2 = d(\mathbf{x}, H^{-1}\mathbf{x}')^2 + d(\mathbf{x}', H\mathbf{x})^2$ (Symmetric transfer error)

Gold standard error

- ▶ Actually, we can consider \mathbf{x} and \mathbf{x}' as noisy observations of ground truth positions $\hat{\mathbf{x}}$ and $\hat{\mathbf{x}}' = H\hat{\mathbf{x}}$.



$$\epsilon(H, \hat{\mathbf{x}}) = d(\mathbf{x}, \hat{\mathbf{x}})^2 + d(\mathbf{x}', H\hat{\mathbf{x}})^2$$

- ▶ **Problem:** this has a lot of parameters: $H, \{\hat{\mathbf{x}}_i\}_{i=1\dots n}$

Sampson error

- ▶ A method that linearizes the dependency on $\hat{\mathbf{x}}$ in the gold standard error so as to eliminate these unknowns.

$$0 = \epsilon(H, \hat{\mathbf{x}}) = \epsilon(H, \mathbf{x}) + J(\hat{\mathbf{x}} - \mathbf{x}) \text{ with } J = \frac{\partial \epsilon}{\partial \mathbf{x}}(H, \mathbf{x})$$

- ▶ Find $\hat{\mathbf{x}}$ minimizing $\|\mathbf{x} - \hat{\mathbf{x}}\|^2$ subject to $J(\mathbf{x} - \hat{\mathbf{x}}) = \epsilon$
- ▶ **Solution:** $\mathbf{x} - \hat{\mathbf{x}} = J^T(JJ^T)^{-1}\epsilon$ and thus:


$$\|\mathbf{x} - \hat{\mathbf{x}}\|^2 = \epsilon^T(JJ^T)^{-1}\epsilon \quad (1)$$

- ▶ Here, $\epsilon_i = A_i h = \mathbf{x}'_i \times (H\mathbf{x}_i)$ is a 3-vector.
- ▶ For each i , there are 4 variables $(\mathbf{x}_i, \mathbf{x}'_i)$, so J is 3×4 .
- ▶ This is almost the algebraic error $\epsilon^T \epsilon$ but with adapted scalar product.
- ▶ The resolution, through iterative method, must be initialized with the algebraic minimization.

Applying homography to image

Two methods:

1. **push** pixels to transformed image and round to the nearest pixel center.
2. **pull** pixels from original image by interpolation.



Contents

Pinhole camera model

Projective geometry

Homographies

Panorama

Internal calibration

Optimization techniques

Conclusion

Practical session


Camera calibration by resection

[R.Y. Tsai, *An efficient and accurate camera calibration technique for 3D machine vision*, CVPR'86] We estimate the camera internal parameters from a known rig, composed of 3D points whose coordinates are known.

- ▶ We have points \mathbf{X}_i and their projection \mathbf{x}_i in an image.
- ▶ In homogeneous coordinates: $\mathbf{x}_i = P\mathbf{X}_i$ or the 3 equations (but only 2 of them are independent)

$$\mathbf{x}_i \times (P\mathbf{X}_i) = 0$$

- ▶ Linear system in unknown P . There are 12 parameters in P , we need 6 points in general (actually only 5.5).
- ▶ Decomposition of P allows finding K .



Restriction: The 6 points cannot be on a plane, otherwise we have a degenerate situation; in that case, 4 points define the homography and the two extra points yield no additional constraint.

Calibration with planar rig

[Z. Zhang *A flexible new technique for camera calibration* 2000]

- ▶ **Problem:** One picture is not enough to find K .
- ▶ **Solution:** Several snapshots are used.
- ▶ For each one, we determine the homography H between the rig and the image.
- ▶ The homography being computed with an arbitrary multiplicative factor, we write

$$\lambda H = K \begin{pmatrix} R_1 & R_2 & T \end{pmatrix}$$

- ▶ We rewrite:

$$\lambda K^{-1} H = \lambda \begin{pmatrix} K^{-1} H_1 & K^{-1} H_2 & K^{-1} H_3 \end{pmatrix} = \begin{pmatrix} R_1 & R_2 & T \end{pmatrix}$$

- ▶ 2 equations expressing orthonormality of R_1 and R_2 :

$$H_1^T (K^{-T} K^{-1}) H_1 = H_2^T (K^{-T} K^{-1}) H_2$$

$$H_1^T (K^{-T} K^{-1}) H_2 = 0$$

- ▶ With 3 views, we have 6 equations for the 5 parameters of $K^{-T} K^{-1}$; then Cholesky decomposition.

The problem of geometric distortion

- ▶ At small or moderate focal length, we cannot ignore the geometric distortion due to lens curvature, especially away from image center.
- ▶ This is observable in the non-straightness of certain lines:



Photo: 5600×3700 pixels



Deviation of 30 pixels

- ▶ The classical model of distortion is radial polynomial:

$$\begin{pmatrix} x_d \\ y_d \end{pmatrix} - \begin{pmatrix} d_x \\ d_y \end{pmatrix} = (1 + a_1 r^2 + a_2 r^4 + \dots) \begin{pmatrix} x - d_x \\ y - d_y \end{pmatrix}$$

Estimation of geometric distortion

- ▶ If we integrate distortion coefficients as unknowns, there is no more closed formula estimating K .
- ▶ We have a non-linear minimization problem, which can be solved by an iterative method.
- ▶ To initialize the minimization, we assume no distortion ($a_1 = a_2 = 0$) and estimate K with the previous linear procedure.

Contents

Pinhole camera model

Projective geometry

Homographies

Panorama

Internal calibration

Optimization techniques

Conclusion

Practical session

Linear least squares problem

- ▶ For example, when we have more than 4 point correspondences in homography estimation:

$$A_{m \times 8} h = B_m \quad m \geq 8$$

- ▶ In the case of an overdetermined linear system, we minimize

$$\epsilon(\mathbf{X}) = \|A\mathbf{X} - B\|^2 = \|f(\mathbf{X})\|^2$$

- ▶ The gradient of ϵ can be easily computed:

$$\nabla \epsilon(\mathbf{X}) = 2(A^T A \mathbf{X} - A^T B)$$

- ▶ The solution is obtained by equating the gradient to 0:

$$\mathbf{X} = (A^T A)^{-1} A^T B$$

- ▶ **Remark 1:** this is correct only if $A^T A$ is invertible, that is A has full rank.
- ▶ **Remark 2:** if A is square, it is the standard solution $\mathbf{X} = A^{-1} B$
- ▶ **Remark 3:** $A^{(-1)} = (A^T A)^{-1} A^T$ is called the pseudo-inverse of A , because $A^{(-1)} A = I_n$.

Non-linear least squares problem

- ▶ We would like to solve as best we can $f(\mathbf{X}) = 0$ with f non-linear. We thus minimize

$$\epsilon(\mathbf{X}) = \|f(\mathbf{X})\|^2$$

- ▶ Let us compute the gradient of ϵ :

$$\nabla\epsilon(\mathbf{X}) = 2J^T f(\mathbf{X}) \text{ with } J_{ij} = \frac{\partial f_i}{\partial x_j}$$

- ▶ Gradient descent: we iterate until convergence

$$\Delta\mathbf{X} = -\alpha J^T f(\mathbf{X}), \alpha > 0$$

- ▶ When we are close to the minimum, a faster convergence is obtained by Newton's method:

$$\epsilon(\mathbf{X}_0) \sim \epsilon(\mathbf{X}) + \nabla\epsilon(\mathbf{X})^T(\Delta\mathbf{X}) + (\Delta\mathbf{X})^T(\nabla^2\epsilon)(\Delta\mathbf{X})$$

and minimum is for $\Delta\mathbf{X} = -(\nabla^2\epsilon)^{-1}\nabla\epsilon$

Levenberg-Marquardt algorithm

- ▶ This is a mix of gradient descent and quasi-Newton method (*quasi* since we do not compute explicitly the Hessian matrix, but approximate it).
- ▶ The gradient of ϵ is

$$\nabla\epsilon(\mathbf{X}) = 2J^T f(\mathbf{X})$$

so the Hessian matrix of ϵ is composed of sums of two terms:

1. Product of first derivatives of f .
 2. Product of f and second derivatives of f .
- ▶ The idea is to ignore the second terms, as they should be small when we are close to the minimum ($f \sim 0$). The Hessian is thus approximated by

$$H = 2J^T J$$

- ▶ Levenberg-Marquardt iteration:

$$\Delta\mathbf{X} = -(J^T J + \lambda I)^{-1} J^T f(\mathbf{X}), \lambda > 0$$

Levenberg-Marquardt algorithm

- ▶ Principle: gradient descent when we are far from the solution (λ large) and Newton's step when we are close (λ small).

1. Start from initial \mathbf{X} and $\lambda = 10^{-3}$.
2. Compute

$$\Delta \mathbf{X} = -(J^T J + \lambda I)^{-1} J^T f(\mathbf{X}), \lambda > 0$$

3. Compare $\epsilon(\mathbf{X} + \Delta \mathbf{X})$ and $\epsilon(\mathbf{X})$:
 - 3a If $\epsilon(\mathbf{X} + \Delta \mathbf{X}) \sim \epsilon(\mathbf{X})$, finish.
 - 3b If $\epsilon(\mathbf{X} + \Delta \mathbf{X}) < \epsilon(\mathbf{X})$,


$$\mathbf{X} \leftarrow \mathbf{X} + \Delta \mathbf{X} \quad \lambda \leftarrow \lambda/10$$

$$3c \text{ If } \epsilon(\mathbf{X} + \Delta \mathbf{X}) > \epsilon(\mathbf{X}), \lambda \leftarrow 10\lambda$$


4. Go to step 2.

Example of distortion correction

Results of Zhang:




Snapshot 1




Snapshot 2

Example of distortion correction

Results of Zhang:



Corrected image 1



Corrected image 2

Contents

Pinhole camera model

Projective geometry

Homographies

Panorama

Internal calibration

Optimization techniques

Conclusion

Practical session

Conclusion

- ▶ Camera matrix K (3×3) depends only on internal parameters of the camera.
- ▶ Projection matrix P (3×4) depends on K and position/orientation.
- ▶ Homogeneous coordinates are convenient as they linearize the equations.
- ▶ A homography between two images arises when the observed scene is flat or the principal point is fixed.
- ▶ 4 or more correspondences are enough to estimate a homography (in general)

Contents

Pinhole camera model

Projective geometry

Homographies

Panorama

Internal calibration

Optimization techniques

Conclusion

Practical session

Practical session: panorama construction

Objective: the user clicks 4 or more corresponding points in left and right images. After a right button click, the program computes the homography and shows the resulting panorama in a new window.

- ▶ Install Imagine++
(<http://imagine.enpc.fr/~monasse/Imagine++/>) on your machine.
- ▶ Let the user click the matching points.
- ▶ Build the linear system to solve $Ax = b$; use `linSolve` (in `LinAlg`) to find the solution x .
- ▶ Compute the bounding box of the panorama.
- ▶ Stitch the images: on overlapping area, take the average of colors at corresponding pixels in both images.